

TimerMeter: Quantifying Timer Method Accuracy and Invocation Cost

Michael.Kuperberg@kit.edu

March 24th, 2010




Forschungszentrum Karlsruhe
in der Helmholtz-Gemeinschaft



Universität Karlsruhe (TH)
Research University · founded 1825

Java Users Group Karlsruhe, Lightning talks



What do you use for timing in Java?

perf4j
getThreadCPUTime
currentTimeMillis
Timestamp
nanoTime
Time
Date

■ What can you conclude from the following?

```
start = System.nanoTime();  
yourMethodToBeBenchmarked();  
duration System.nanoTime()-start; //e.g. 1955 ns
```

really 1955 ns?

■ `java.lang.System.nanoTime()`


- from official javadocs: *“nanosecond precision, but not necessarily nanosecond accuracy”*
- no API-provided means to obtain precision/accuracy
- anecdotal evidence on the WWW, different results:
e.g. accuracy „a few hundred nanoseconds“

Overview

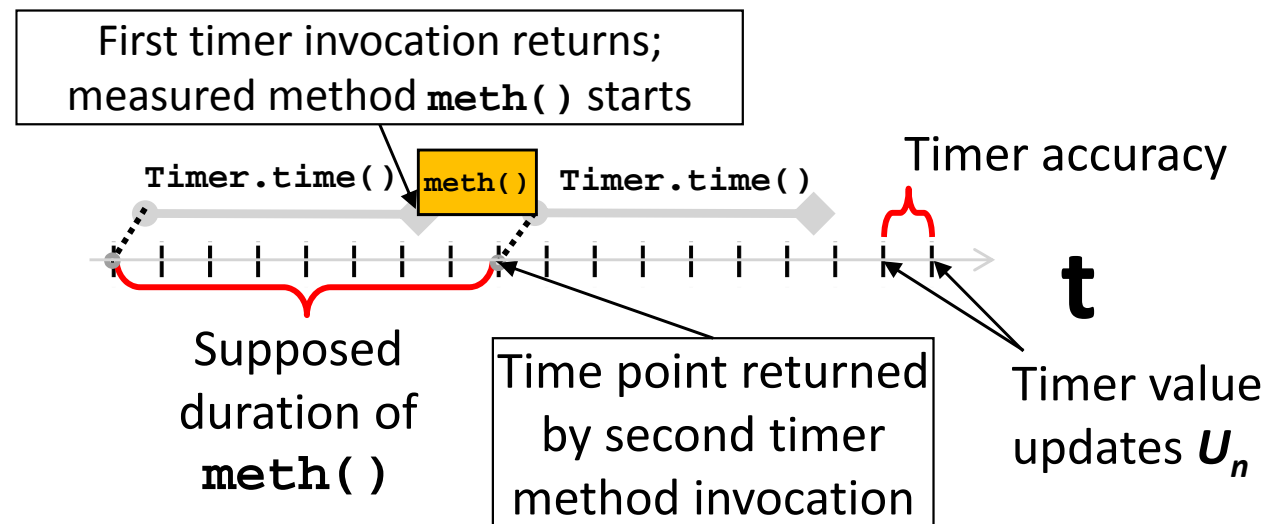
- ✓ **Motivation**
 - **Foundations**
 - **Requirements**
 - **Main Idea of TimerMeter**
 - **Evaluation**
 - **Conclusion**
 - **(Related Work)**

TimerMeter: Foundations (1)

Effect of accuracy on measurements

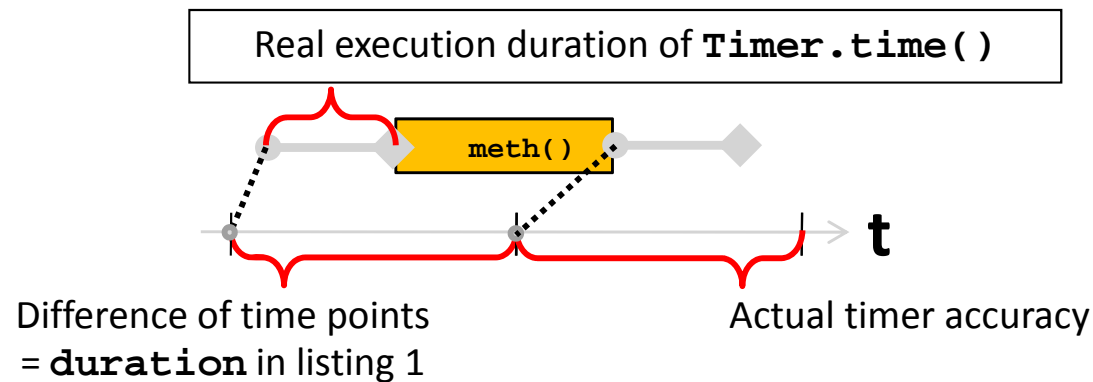


Case 1: accuracy < execution duration



TimerMeter: Foundations (2)

■ Case 2: $\text{accuracy} \geq \text{execution duration}$



- Thus, we need to know both accuracy and execution duration!
- It's hard to disentangle the measurement overhead from what's being measured


TimerMeter: Requirements

- Context: timers for fine-grained benchmarks and measurements (→ less than a microsecond)
- Usual approach is „take the best available timer“
 - Java: many timer methods available, incl. 3rd-party
 - The choice is often not clear, not justifiable – or wrong
- Thus: **we need to know timer accuracy / resolution**
 - HW-specific and OS-specific
 - they contribute to the measured timing values
 - they impact the statistical quality of results

Solution: next slides!

TimerMeter: The Main Idea

- Central idea of TimerMeter: gradually and slowly increase work between timer invocations
 - so that the measured interval increases by an accuracy at some point (cf. paper)
 - works for $\underline{\text{Accuracy}} < \underline{\text{InvocationCost}}$ and $A \geq IC$
- Implementation is more complex
 - Clusters due to rounding and truncation etc.



TimerMeter: Evaluation: Results

★ ≡ „calculated from frequency“, ◇ ≡ „invocation cost measured using nanoTime()“


		Platform P1			
		Linux 2.6.25 JDK 1.6.0_07		Win XP JDK 1.6.0_07	
Timer method / Counter	Precision unit	Accuracy	Cost	Accuracy	Cost
rdtsc	CPU Cycle	2	130	2	106
QueryPerformanceCounter	279.4 ns★	n/a	n/a	1	6
nanoTime	ns	70	978	279	1676
highResCounter (Linux)	1000 ns★	1	2	-	-
highResCounter (Windows)	279.4 ns★	-	-	1	7
currentTimeMillis ◇	ms	1	0.004 ◇	15	0.0002 ◇
getCurrentThreadCpuTime ◇	ns	15·10 ⁶	786 ◇	15.6·10 ⁶	27 ◇
JETM	ns	70	978	279	1676

- Note the difference of OSes for nanoTime ()
- Note the accuracy of getCurrentThreadCpuTime
- JETM: very similar to nanoTime (): accuracies correspond to (rounded) value of 1 HPET counter tick

The Larger Context:

<http://www.palladio-approach.net>

Model-based Architecture Performance Evaluation



- Research on software architectures (KIT and FZI): benchmarking, reverse engineering, reliability, etc.
- Development using Eclipse: GEF, GMF, GEF etc.
- Student theses (master, bachelor) and jobs available

TimerMeter: Conclusions

- TimerMeter: a novel, easily portable algorithm
 - For transparently quantifying accuracy and invocation cost of timer methods on your platform
 - Available under EPL. Details, docs etc.:
<http://bit.ly/TimerMeter> or <https://sdqweb.ipd.kit.edu/wiki/TimerMeter>
- The accuracy of a timer method can differ by 10x or more depending on OS (e.g. `currentTimeMillis`)
 - `nanoTime()`: accuracy „only“ 70 ns up to 279 ns
 - `nanoTime()`: invocation cost: 978 ns up to 1676 ns